

Andrés Paredes

Senior Data Scientist

andres.paredes@email.com · +54 9 11 5533-9912 · Buenos Aires, Argentina · kaggle.com/andresparedes

PROFILE SUMMARY

Data scientist with 6 years of experience delivering end-to-end ML solutions from problem framing to production deployment. Expertise in predictive modelling, NLP, and recommendation systems. Published 2 papers in applied ML conferences. Strong communicator: translates complex results for non-technical stakeholders.

WORK EXPERIENCE

Senior Data Scientist

[Naranja X](#)

Jun 2021 – Present · Buenos Aires (hybrid)

- Built a real-time credit risk model (XGBoost + SHAP) that reduced default rate by 1.8 pp.
- Led development of a product recommendation engine (collaborative filtering + content-based); lifted CTR by 22%.
- Set up MLflow experiment tracking and model registry for 4 data science squads.
- Designed A/B testing framework used across 15+ model experiments per quarter.

Data Scientist

[Satellogic](#)

Jan 2019 – May 2021 · Buenos Aires, Argentina

- Developed CNN-based image segmentation models for satellite imagery analysis (TensorFlow).
- Built automated data pipelines with Apache Airflow and dbt; reduced processing time by 50%.
- Collaborated with the engineering team to serve models via a FastAPI microservice.

Junior Data Analyst

[Consultora Analytics BA](#)

Mar 2018 – Dec 2018 · Buenos Aires, Argentina

- Built dashboards and statistical reports for retail and FMCG clients using Python and Tableau.

EDUCATION

M.Sc. Data Science

[Universidad de San Andrés](#)

2017 – 2018

B.Sc. Mathematics

[Universidad de Buenos Aires \(UBA\)](#)

2012 – 2016

TECHNICAL SKILLS

Languages	Python (expert), R (intermediate), SQL
ML / DL	scikit-learn, XGBoost, LightGBM, TensorFlow, PyTorch, Hugging Face
Data Engineering	Apache Airflow, dbt, pandas, PySpark (basic)
Experimentation	A/B testing, causal inference, MLflow, Weights & Biases
Databases	PostgreSQL, BigQuery, Redshift
Practices	CRISP-DM, reproducible research, model interpretability (SHAP)

LANGUAGES

Spanish	Native
---------	--------

PROJECTS

ChurnPredictor OSS

Open-source customer churn prediction toolkit built with scikit-learn, including automated feature engineering and threshold optimization. 210 GitHub stars.